

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df=pd.read_csv(r"C:\Users\Niranjan\Downloads\used_cars_data.csv")
df
```

```
Out[2]:
```

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	5
...	...	...	...	...	...	...	...	...
7248	7248	Volkswagen Vento Diesel Trendline	Hyderabad	2011	89411	Diesel	Manual	
7249	7249	Volkswagen Polo GT TSI	Mumbai	2015	59000	Petrol	Automatic	
7250	7250	Nissan Micra Diesel XV	Kolkata	2012	28000	Diesel	Manual	
7251	7251	Volkswagen Polo GT TSI	Pune	2013	52262	Petrol	Automatic	
7252	7252	Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan...	Kochi	2014	72443	Diesel	Automatic	

7253 rows × 14 columns

In [3]: `df.head()`

Out[3]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	Firs
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	Firs
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	Firs
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	Firs
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Secon

In [4]: `df.tail()`

Out[4]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type
7248	7248	Volkswagen Vento Diesel Trendline	Hyderabad	2011	89411	Diesel	Manual	
7249	7249	Volkswagen Polo GT TSI	Mumbai	2015	59000	Petrol	Automatic	
7250	7250	Nissan Micra Diesel XV	Kolkata	2012	28000	Diesel	Manual	
7251	7251	Volkswagen Polo GT TSI	Pune	2013	52262	Petrol	Automatic	
7252	7252	Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan...	Kochi	2014	72443	Diesel	Automatic	

In [5]: `df.describe()`

Out[5]:

	S.No.	Year	Kilometers_Driven	Seats	Price
<b>count</b>	7253.000000	7253.000000	7.253000e+03	7200.000000	6019.000000
<b>mean</b>	3626.000000	2013.365366	5.869906e+04	5.279722	9.479468
<b>std</b>	2093.905084	3.254421	8.442772e+04	0.811660	11.187917
<b>min</b>	0.000000	1996.000000	1.710000e+02	0.000000	0.440000
<b>25%</b>	1813.000000	2011.000000	3.400000e+04	5.000000	3.500000
<b>50%</b>	3626.000000	2014.000000	5.341600e+04	5.000000	5.640000
<b>75%</b>	5439.000000	2016.000000	7.300000e+04	5.000000	9.950000
<b>max</b>	7252.000000	2019.000000	6.500000e+06	10.000000	160.000000

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   S.No.                 7253 non-null  int64
1   Name                  7253 non-null  object
2   Location              7253 non-null  object
3   Year                  7253 non-null  int64
4   Kilometers_Driven     7253 non-null  int64
5   Fuel_Type             7253 non-null  object
6   Transmission          7253 non-null  object
7   Owner_Type            7253 non-null  object
8   Mileage               7251 non-null  object
9   Engine                7207 non-null  object
10  Power                 7207 non-null  object
11  Seats                 7200 non-null  float64
12  New_Price             1006 non-null  object
13  Price                 6019 non-null  float64
dtypes: float64(2), int64(3), object(9)
memory usage: 793.4+ KB
```

In [8]: `df.shape`

Out[8]: (7253, 14)

```
In [9]: df.isna().any()
```

```
Out[9]: S.No.                False
        Name                False
        Location            False
        Year                False
        Kilometers_Driven    False
        Fuel_Type            False
        Transmission         False
        Owner_Type           False
        Mileage              True
        Engine               True
        Power                True
        Seats                True
        New_Price            True
        Price                True
        dtype: bool
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: S.No.                0
        Name                0
        Location            0
        Year                0
        Kilometers_Driven    0
        Fuel_Type            0
        Transmission         0
        Owner_Type           0
        Mileage              2
        Engine              46
        Power               46
        Seats               53
        New_Price           6247
        Price              1234
        dtype: int64
```

```
In [11]: df.fillna(value=0,inplace=True)
```

```
In [12]: df.isnull().sum()
```

```
Out[12]: S.No.          0
         Name          0
         Location      0
         Year          0
         Kilometers_Driven  0
         Fuel_Type      0
         Transmission  0
         Owner_Type     0
         Mileage        0
         Engine         0
         Power          0
         Seats          0
         New_Price      0
         Price          0
         dtype: int64
```

```
In [13]: df=df[['Year','Price']]
         df.columns=['yr','pr']
```


```
In [14]: df.head(10)
```

```
Out[14]:
```

	yr	pr
0	2010	1.75
1	2015	12.50
2	2011	4.50
3	2012	6.00
4	2013	17.74
5	2012	2.35
6	2013	3.50
7	2016	17.50
8	2013	5.20
9	2012	1.95

```
In [15]: sns.lmplot(x="yr",y="pr",data=df,order=2,ci=None)
```

```
Out[15]: <seaborn.axisgrid.FacetGrid at 0x2215b0c95d0>
```



```
In [16]: df.describe()
```

```
Out[16]:
```

	yr	pr
count	7253.000000	7253.000000
mean	2013.365366	7.866665
std	3.254421	10.796286
min	1996.000000	0.000000
25%	2011.000000	2.290000
50%	2014.000000	4.650000
75%	2016.000000	8.400000
max	2019.000000	160.000000

```
In [17]: df.fillna(method='ffill',inplace=True)
```

C:\Users\Niranjan\AppData\Local\Temp\ipykernel\_10700\4116506308.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
df.fillna(method='ffill',inplace=True)

```
In [18]: x=np.array(df['yr']).reshape(-1,1)
y=np.array(df['pr']).reshape(-1,1)
```

```
In [19]: df.dropna(inplace=True)
```


C:\Users\Niranjan\AppData\Local\Temp\ipykernel\_10700\1379821321.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
df.dropna(inplace=True)

```
In [20]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.0601604846904914


```
In [21]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='b')
plt.scatter(X_test,y_pred,color='k')
plt.show()
```






```
In [22]: df500=df[:500]  
sns.lmplot(x="yr",y="pr",data=df500,order=1,ci=None)
```

Out[22]: <seaborn.axisgrid.FacetGrid at 0x2215dd36250>



```
In [23]: df500.fillna(method='ffill',inplace=True)
X=np.array(df500['yr']).reshape(-1,1)
y=np.array(df500['pr']).reshape(-1,1)
df500.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr=LinearRegression()
regr.fit(X_train,y_train)
print("Regression:",regr.score(X_test,y_test))
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
```

Regression: 0.16168775787920153



```
In [24]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.16168775787920153

In [ ]: